


# Binary Search Trees, Heaps, SkipLists



1332 Recitation: Week of May 25th


# Trees Intro

- ADT where nodes have references to child nodes
- Recursive in structure (each node is the “root” of a smaller subtree)
- An “internal” node is one that has children
- An “external” node or “leaf” does not have children
- *Binary Tree*
  - Each node has  $\leq 2$  children
- *Depth*: how many edges/references away from the root a node is
- *Height*:  $\max(\text{left child height}, \text{right child height}) + 1$ , height of a null is -1
  - This means leaves have height 0




# Tree Shape Properties


- *Full*: all internal nodes (nodes with children) have the *maximum* number of children
- *Balanced*: the height of the tree is  $O(\log n)$ 
  - Becomes a more important concept later, will be redefined later
- *Complete*: all levels are completely filled in, except the last level, which is filled from left to right
  - A complete tree is always balanced
- *Degenerate*: each node has appr. 1 child, height of tree is  $O(n)$




Full tree



Complete tree



Full & complete tree



Degenerate tree

# Binary Search Trees!

- A binary tree with an *order property*, which requires data to be Comparable
  - All data to the left of a node is  $< \text{node.data}$
  - All data to the right of a node is  $> \text{node.data}$
  - No duplicates in this class
- To search:
  - Start with  $\text{current} = \text{root}$
  - If  $\text{current} == \text{null}$ , data not found
  - If  $\text{data} < \text{current}$ , recursive call on  $\text{current.left}$
  - If  $\text{data} > \text{current}$ , recursive call on  $\text{current.right}$
- If the BST is balanced, search is  $O(\log n)$ , since there are  $O(\log n)$  levels
- If the BST is imbalanced/degenerate, search is  $O(n)$ , since there are  $O(n)$  levels

# Pointer Reinforcement!

- A way to recursively modify the tree without having to “look ahead”
  - Remember LL? We had to stop on the node *before* the index
  - Look-ahead is very ugly and unclean for BST operations (please don't do it)
  - If you learn pointer reinforcement now, your HW 5 will be mostly copy-paste
- How does it work?
  - Method Signature: **BSTNode** foo(BSTNode curr, ...);
    - Takes in a node, *returns* a node
  - When making a recursive call to the left, set **curr.left = foo(curr.left, ...)**
  - The node you return at any given recursive call is the **new root** of its subtree
    - This means **return curr**; makes no change to the structure (it *reinforces* the pointer), and returning a different node changes the structure
  - In the public method, you have to reinforce the root: **root = foo(root, ...);**

# BST Add

- We always add to the very bottom of the tree, creating a new leaf node
- Basically just a search, but with pointer reinforcement:
  - If  $\text{data} < \text{curr}$ , **`curr.left = addR(curr.left)`**, then return curr
  - If  $\text{data} > \text{curr}$ , **`curr.right = addR(curr.right)`**, then return curr
  - If  $\text{data} == \text{curr}$ , then we found a duplicate, so just return curr
  - When  $\text{curr} == \text{null}$ , we have found the spot where the new node should go
    - *Return a new node with the data*
      - Note that **`return curr`** would return null, which was already the parent's child
      - But **`return newNode`** changes the parent's reference from null to newNode
- Example!

# BST Remove

- See “Snowcitation” video from Spring 2019 for implementation details:  
*[bit.ly/2m4iWSV](https://bit.ly/2m4iWSV)*, also available on Piazza on Guide to Success & Resources in CS 1332
- Traverse to find data, similar to add. Once found, there are 3 cases
  - 0 children: just return null
  - 1 child: return the **non-null** child
  - 2 children: replace curr’s data with the predecessor/successor and remove the pred/succ
    - 5, 10, 25, 39, 50, 51, 99, 101; 25 is 39’s predecessor, 50 is 39’s successor
    - To find and remove the pred/succ, have another pointer-reinforced helper:
      - The pred. can be found by traversing left once, then right until curr.right is null
      - The succ. can be found by traversing right once, then left until curr.left is null
      - Once found, store the data in a dummy node, and return the pred/succ non-null child to remove the pred/succ

# BST Traversals

- Two categories: breadth-first and depth-first
- *Breadth-first*: levelorder(), requires a queue and a while-loop
- *Depth-first*: preorder(), inorder(), and postorder(), requires recursion

levelorder():

```
if root not null, put root in queue
while (queue not empty):
    node <- queue.dequeue()
    visit node
    enqueue all node's non-null children
```

preorder(curr):

```
if curr not null:
    visit curr
    preorder(curr.left)
    preorder(curr.right)
```

For inorder(), go left, then visit, then go right

For postorder(), go left, then go right, then visit



# Heaps!

- A binary tree backed by an array
- Two types, both only care about one element at any given point:
  - MinHeap always keeps track of the smallest value
  - MaxHeap always keeps track of the greatest value
- Properties:
  - SHAPE: Always **complete** (all levels filled except last level, which is filled from left to right)
  - ORDER: Each element is less/greater than both its children
  - Note that it is easier to “fix” a broken order property than a broken shape property, so we will prioritize keeping completeness when doing operations

# Array-backed tree? How?

- The completeness (lack of gaps) allows us to easily use an array to model their structure
- `arr[0]` is left empty for the sake of index arithmetic
- `arr[1]` is the root, fill indices in sequentially in level-order
- For any element at index  $i$ :
  - Its left child is at index  $(i * 2)$
  - Its right child is at index  $(i * 2 + 1)$
  - Its parent is at index  $(i / 2)$

# Heap Add

- Put the new element at the end of the array, then upheap()
- upheap()
  - Compare element,  $\text{arr}[i]$ , to its parent -- if order property is not satisfied, then swap the element with its parent (swap  $\text{arr}[i]$  with  $\text{arr}[i/2]$ )
  - If a swap occurred, then repeat on the parent index (set  $i = i/2$  and repeat)
  - As soon as order property is satisfied, no more comparisons should occur.
- $O(\log n)$  -> max of 1 swap per level, and since the tree is complete, it is always balanced

# Heap Remove



- Remember, we only care about the min/max at the root! This is the only element we will ever remove.
- Store the root (`arr[1]`) to return, and then replace that index with the very last element
- Now `downheap()`:
  - Compare the element to its children -- if the order property is not satisfied, then a swap must occur. Swap with the **smaller child in a MinHeap** and the **larger child in a MaxHeap**
  - Repeat until the order property is satisfied.
- $O(\log n)$  -> max of 1 swap per level, and since the tree is complete, it is always balanced

# BuildHeap Algorithm

- When constructing a heap, you could add one-by-one
  - Add costs  $O(\log n)$ , repeated  $n$  times  $\rightarrow O(n \log n)$
- We can actually do better, using the BuildHeap algorithm
  - Put all the elements in the backing array in their current order
  - Start at the last internal node (i.e. the node closest to the back that still has children, located at index  $(\text{size} / 2)$ )
  - for  $(i = \text{size} / 2; i > 0; i--)$  {  $\text{downheap}(i);$  }
- Due to some sequence/series math, this algorithm is actually  $O(n)$ , much better than  $O(n \log n)$


# SkipLists!

- A different data structure that uses probability to decide its structure
- Basically a bunch of DLLs stacked on top of each other
  - Nodes have left/right, like prev/next in a DLL
  - Also have up/down to connect up the DLLs
  - `node.up.data == node.down.data == node.data`
- Each level contains a subset of the data of the data below it
- We choose what data goes on what level using coin flips



# SkipList Search

- Top-left infinity node is our entry point to the list
  - Check curr.right:
    - if less than data, move curr = curr.right
    - else, drop down a level and repeat
  - If we hit null, the data isn't found
- Ideally, we have  $n$  data on the bottom level, then  $n/2$  on the next level, then  $n/4$ , etc., which makes us able to skip half the data at each level, resulting in  $O(\log n)$  time



# SkipList Add

- When adding, we search for the spot where the element *should* be
- All elements go on the bottom level
- The number of times we promote an element depends on the number of times we flip a coin before hitting a “tails”
  - Heads = promote
  - Tails = stop promoting and place the new nodes
- If there aren't enough levels, then add a new one
  - We typically cap the number of levels at appr.  $\log(n)$ , so that the worst-case space complexity is  $O(n \log n)$



# SkipList Remove

- One of the few structures where removing is easier than adding
- Simply traverse to the data, and remove all instances of the data by using the up/down pointers in the nodes